

Alpaca Options Agent

Deterministic Risk & Statistical Validation

An architecture that refuses bad trades.

1. The Determinism Problem

The central problem in agentic trading is not intelligence — it is reproducibility.

We ran the agent 40 times on the same prompt at temperature 0 with a fixed seed.

- **High Divergence:** A significant number of replays diverged despite fixed parameters.
- **Tool Flips:** Replays changed the *tool called entirely*, not just arguments. One original call flipped from reading an option chain to attempting an order.

Conclusion: The model is not reproducible under conditions meant to guarantee it. Authority must sit in deterministic Python, not the LLM.

2. Adversarial Validation of the Risk Gate

We ran the agent's risk gate against itself to prove it works under adversarial conditions. Two critical holes were found and closed:

1. **Ratio Quantity Exploit:** A buy-1/sell-2 payload was previously priced as a 1:1 vertical (second short leg treated as zero-cost). The gate now correctly parses and stops this.
2. **Hallucinated Strikes:** The gate previously approved a hallucinated OCC symbol with a non-existent strike. It now validates every leg against the live chain.

Both were found by the agent attacking itself.

3. Statistical Validation Gate

We evaluated 21 distinct strategy/symbol pairs (SPY, QQQ, IWM $\times$ 7 structures), producing 24 total validation records.

- **0 pairs cleared for live trading.**
- The closest: SPY vertical credit spread (78% win rate, Sharpe 1.67). Refused because the Sharpe CI lower bound remained negative (-0.32).
- **Moving-block bootstrap:** Correcting from an i.i.d. bootstrap to a moving-block model cut the false-positive rate from **10.8%** to a calibrated **2.5%**.

The gate governs claims of edge. The risk layer governs structural safety.

4. The Proposer / Critic Architecture

The LLM never computes a number that reaches an order.

- **Proposer Agent:** Suggests regime and parameters.
- **Critic Agent:** Evaluates the proposal against our validation graveyard.
- **Featherless Arbiter:** A third open-weights model arbitrates disagreements.

Live Veto: The Critic agent actively vetoed a proposed MSFT cash-secured put for citing no backtest evidence on a symbol that never cleared the gate.

Cost: \$0.022 per Proposer/Critic cycle.

5. Deterministic Risk Gate & Execution

Every order must pass the non-LLM Python risk gate before routing to the Alpaca MCP server.

- We implemented a batch model for structure-level capital at risk.
- **Old per-leg model:** Refused \$75,500 capital-at-risk for a SPY spread.
- **New batch model:** Correctly calculated \$[FINAL_BATCH_RISK] risk and approved the trade.

Demonstration Trade: We executed **one** unvalidated SPY credit spread pinned to the Friday Sept 4th expiry, explicitly bypassing validation to demonstrate the execution path. We close Thursday to avoid Friday's NFP event risk.

Thank You

Code public.

Options quotes provided by Alpaca's free-tier Indicative Pricing Feed.